Building a Context Pruning Pipeline for Long-Running Agents	Modern AI agents built on top of large language models (LLMs) are designed to run continuously. As a result, their conversation history keeps growing indefinitely. Passing such an entire history as the LLM's context window is the perfect recipe for prohibitive token costs, latency bottlenecks, and eventual degradation in reasoning. Building a context pruning pipeline can address this issue by dynamically managing recent conversational memory. https://machinelearningmastery.com/building-a-context-pruning-pipeline-for-long-running-agents/	CDC, BDA, AIML, CSE